



deeplearning.ai

Setting up your
ML application

Train/dev/test
sets

Applied ML is a highly iterative process


layers

hidden units

learning rates

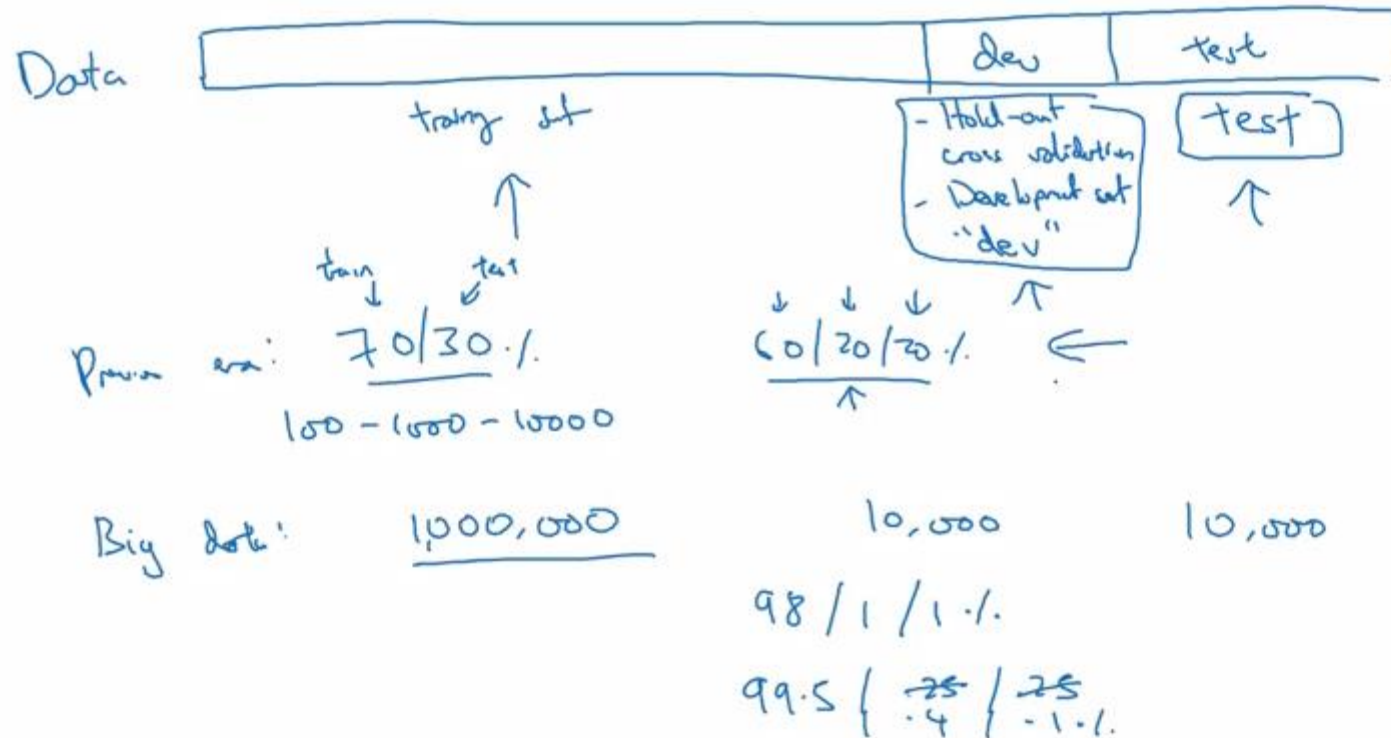
activation functions

...



NLP, Vision, Speech, Structured Data
└─┬──────────┘
 | |
 └─┬─┐ |
 | | |
 └─┴─┘ |
 |
 └─┬─┐
 | |
 └─┴─┘
 Search Security Logistic ...

Train/dev/test sets



Mismatched train/test distribution

Certs

↓
Training set:

Cat pictures from
webpages

↓ ↓
Dev/test sets:

Cat pictures from
users using your app

→ Make sure dev and test come from same distribution.

↓ ↓
train / test
↓ ↓
train / dev
↑

train / test
↓ ↑
→ train / dev

Not having a test set might be okay. (Only dev set.)






deeplearning.ai

Setting up your
ML application

Bias/Variance

Bias and Variance




Bias and Variance

Cat classification



Train set error:	1%	15% ↙	15%	0.5%
Dev set error:	11%	16% ↙	30%	1%
	high variance ↑	high bias ↑↑	high bias & high variance	low bias low variance ↑
Human: ~0%				
Optimal (Bayes) error: ~0%	15%	Blurry images		

High bias and high variance





deeplearning.ai

Setting up your ML application

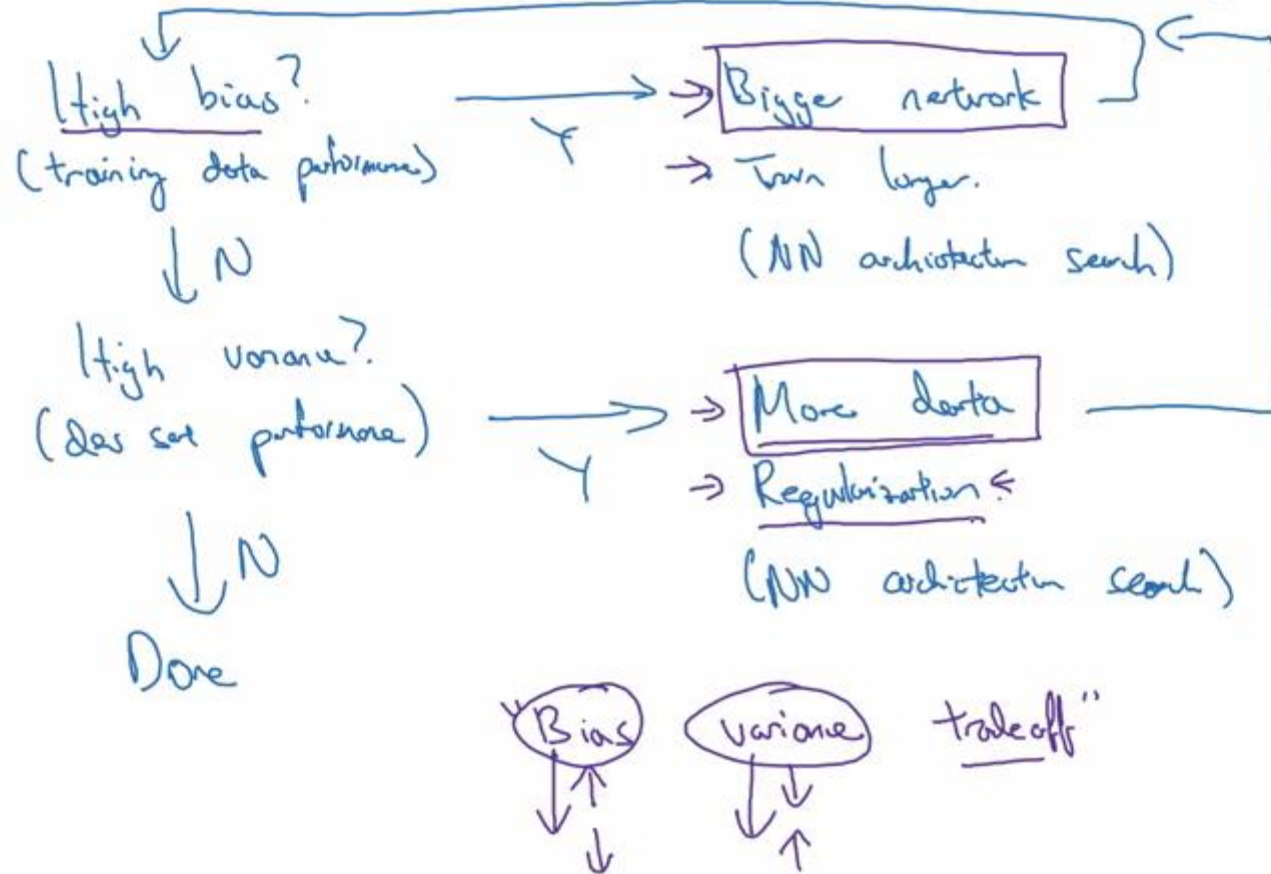
Basic “recipe” for machine learning



0:21 / 6:21



Basic recipe for machine learning





deeplearning.ai

Regularizing your
neural network

Regularization

Logistic regression

$$\min_{w,b} J(w,b)$$

$$\underline{w} \in \mathbb{R}^{n_x}, \underline{b} \in \mathbb{R}$$

λ = regularization parameter
lambda lambda

$$J(w,b) = \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(\hat{y}^{(i)}, y^{(i)})}_{\text{cost}} + \frac{\lambda}{2m} \underbrace{\|w\|_2^2}_{\text{L2 regularization}}$$

~~$$+ \frac{\lambda}{2m} b^2$$~~
 omit

L_2 regularization $\underline{\|w\|_2^2} = \sum_{j=1}^{n_x} w_j^2 = w^T w \leftarrow$

L_1 regularization $\frac{\lambda}{2m} \sum_{j=1}^{n_x} |w_j| = \frac{\lambda}{2m} \|w\|_1$

w will be sparse

Neural network

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(y^{(i)}, \hat{y}^{(i)})}_{\text{loss}} + \underbrace{\frac{\lambda}{2m} \sum_{l=1}^L \|w^{[l]}\|_F^2}_{\text{regularization}}$$

$$\|w^{[l]}\|_F^2 = \sum_{i=1}^{n^{[l-1]}} \sum_{j=1}^{n^{[l]}} (w_{ij}^{[l]})^2$$

"Frobenius norm"

$$\|\cdot\|_2^2 \uparrow$$

$$\|\cdot\|_F^2$$

$$w: \begin{pmatrix} n^{[l]} & n^{[l-1]} \\ \uparrow & \uparrow \end{pmatrix}$$

$$dw^{[l]} = \boxed{(\text{from backprop}) + \frac{\lambda}{m} w^{[l]}}$$

$$\rightarrow w^{[l]} := w^{[l]} - \alpha dw^{[l]}$$

$$\frac{\partial J}{\partial w^{[l]}} = dw^{[l]}$$

"Weight decay"

$$\begin{aligned} w^{[l]} &:= w^{[l]} - \alpha \left[(\text{from backprop}) + \frac{\lambda}{m} w^{[l]} \right] \\ \boxed{\left(1 - \frac{\alpha \lambda}{m}\right) w^{[l]}} &= \boxed{w^{[l]} - \left(\frac{\alpha \lambda}{m}\right) w^{[l]}} - \alpha (\text{from backprop}) \end{aligned}$$




deeplearning.ai

Regularizing your neural network


Why regularization reduces overfitting



How does regularization prevent overfitting?



How does regularization prevent overfitting?



$$g(z) = \tanh(z)$$

$\lambda \uparrow$


$W^{[L]} \downarrow$

$$z^{[L]} = W^{[L]} a^{[L-1]} + b^{[L]}$$



Every layer \approx linear.

$$J(\dots) = \underbrace{\sum_i \mathcal{L}(\hat{y}^{(i)}, y^{(i)})}_{\text{Loss}} + \underbrace{\frac{\lambda}{2} \sum_l \|W^{[l]}\|_F^2}_{\text{Regularization}}$$







deeplearning.ai

Regularizing your neural network

Dropout regularization

Dropout regularization



Implementing dropout ("Inverted dropout")

Illustrate with layer $l=3$. $\text{keep-prob} = \frac{0.8}{x}$ 0.2

$\rightarrow \boxed{d3} = \text{np.random.rand}(a3.\text{shape}[0], a3.\text{shape}[1]) < \text{keep-prob}$

$\underline{a3} = \text{np.multiply}(a3, d3)$ $\# a3 \neq d3$

$\rightarrow \boxed{a3 /= \text{keep-prob}}$ \leftarrow

50 units. \leadsto 10 units shut off

$$z^{[4]} = w^{[4]} \cdot \underline{a^{[3]}} + b^{[4]}$$

\uparrow reduced by 20%.

$$/= \underline{0.8}$$

Test

Making predictions at test time

$$a^{[0]} = X$$

No drop out.

$$z^{[1]} = W^{[1]} \underline{a^{[0]}} + b^{[1]}$$

$$a^{[1]} = g^{[1]}(z^{[1]})$$

$$z^{[2]} = W^{[2]} \underline{a^{[1]}} + b^{[2]}$$

$$a^{[2]} = \dots$$

↓
 \hat{y}

\neq keep-prob





deeplearning.ai

Regularizing your
neural network

Understanding
dropout

Why does drop-out work?

Intuition: Can't rely on any one feature, so have to spread out weights. \leadsto Shrink weights. b_2





deeplearning.ai

Regularizing your
neural network

Other regularization
methods

Data augmentation



4



4

4

4



2:40 / 8:23




Andrew Ng

Early stopping

Orthogonalization.

Optimize cost function J
- Gradient, ...
Not overfit.
- Regularization, ...

$J(w, b)$






deeplearning.ai

Setting up your
optimization problem

Normalizing inputs

Normalizing training sets



Subtract mean:

$$\mu = \frac{1}{n} \sum_{i=1}^m x^{(i)}$$

$$x := x - \mu$$

Normalize variance

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^m x^{(i)} * x^{(i)} \quad \leftarrow \text{element-wise}$$



$$x /= \sigma^2$$

Use same μ σ^2 to normalize test set.

Why normalize inputs?



$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Unnormalized:
 $w_1, x_1: 1 \dots 1000 \leftarrow$
 $w_2, x_2: 0 \dots 1 \leftarrow$
 $-1 \dots 1$



$x_1: 0 \dots 1$
 $x_2: -1 \dots 1$
 $x_3: 1 \dots 2$

Normalized:






deeplearning.ai

Setting up your
optimization problem

Vanishing/exploding
gradients

Vanishing/exploding gradients



$g(z) = z$
 $\hat{y} = w^{(L)} \cdot \dots \cdot w^{(1)} x$
 $b^{(L)} = 0$

$w^{(1)} > I$
 $w^{(2)} < I$

$w^{(1)} = \begin{bmatrix} 1.5 & 0 \\ 0 & 0.5 \end{bmatrix}$

$\hat{y} = w^{(L)} \begin{bmatrix} 1.5^{L-1} & 0 \\ 0 & 0.5^{L-1} \end{bmatrix} x$

$z^{(1)} = w^{(1)} x$
 $a^{(1)} = g(z^{(1)}) = z^{(1)}$
 $a^{(2)} = g(z^{(2)}) = g(w^{(2)} a^{(1)})$

1.5^L
 0.5^L



deeplearning.ai

Setting up your
optimization problem


Weight initialization
for deep networks



0:03 / 6:11



Single neuron example



$$z = w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

large $n \rightarrow$ Smaller w_i

$$\text{Var}(w_i) = \frac{1}{n} \frac{2}{n}$$

$$\underline{w^{[L]}} = n.p. \text{ random} . \underline{\text{randn}}(\text{shape}) * n.p. \text{sqrt}\left(\frac{2}{n^{[L-1]}}\right)$$

ReLU $g^{[L]}(z) = \text{ReLU}(z)$

Other variants:

tanh

$$\frac{1}{n^{[L-1]}}$$

Xavier initialization ↑

$$\sqrt{\frac{2}{n^{[L-1]} + n^{[L]}}}$$

↑



deeplearning.ai


Setting up your
optimization problem

Numerical approximation
of gradients

Checking your derivative computation

云课堂

$$\underline{f(\theta) = \theta^3}$$



$$\frac{f(\theta+\epsilon) - f(\theta-\epsilon)}{2\epsilon} \approx \underline{g(\theta)}$$

$$\frac{(1.01)^3 - (0.99)^3}{2(0.01)} = 3.0001 \approx 3$$

$$g(\theta) = 3\theta^2 = 3$$

approx error: 0.0001

(prev slide: 3.0301, error: 0.03)

$$\left\{ \begin{array}{l} f'(\theta) = \lim_{\epsilon \rightarrow 0} \frac{f(\theta+\epsilon) - f(\theta-\epsilon)}{2\epsilon} \quad \begin{array}{l} O(\epsilon^2) \\ 0.01 \\ 0.0001 \end{array} \quad \left| \quad \frac{f(\theta+\epsilon) - f(\theta)}{\epsilon} \quad \begin{array}{l} \text{error: } O(\epsilon) \\ 0.01 \end{array} \end{array} \right.$$

Andrew Ng



deeplearning.ai

Setting up your
optimization problem

Gradient Checking

Gradient check for a neural network

Take $W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}$ and reshape into a big vector θ .

concentrate

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = J(\theta)$$

Take $dW^{[1]}, db^{[1]}, \dots, dW^{[L]}, db^{[L]}$ and reshape into a big vector $d\theta$.

concentrate

Is $d\theta$ the gradient of $J(\theta)$?

Gradient checking (Grad check)

$$J(\theta) = J(\theta_1, \theta_2, \theta_3, \dots)$$

for each i :

$$\rightarrow \underline{d\theta_{\text{approx}}[i]} = \frac{J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i + \epsilon}, \dots) - J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i - \epsilon}, \dots)}{2\epsilon}$$

$$\approx \underline{d\theta[i]} = \frac{\partial J}{\partial \theta_i} \quad \Bigg| \quad d\theta_{\text{approx}} \approx d\theta$$

Checks

$$\rightarrow \frac{\|d\theta_{\text{approx}} - d\theta\|_2}{\|d\theta_{\text{approx}}\|_2 + \|d\theta\|_2}$$

$\epsilon = 10^{-7}$

$$\approx \frac{10^{-7}}{10^{-5}} - \text{great!} \leftarrow$$

$\rightarrow 10^{-3} - \text{worry.} \leftarrow$



deeplearning.ai

Setting up your
optimization problem

Gradient Checking
implementation notes

Gradient checking implementation notes

- Don't use in training – only to debug

$$\frac{d\theta_{\text{approx}}[i]}{\uparrow \uparrow} \longleftrightarrow \frac{d\theta[i]}{\uparrow}$$

- If algorithm fails grad check, look at components to try to identify bug.

$$\underline{db^{[L]}} \quad \underline{dw^{[L]}}$$

- Remember regularization.

$$\underline{J(\theta)} = \frac{1}{n} \sum_i \ell(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2n} \sum_l \|w^{[l]}\|_F^2$$

$d\theta = \text{gradient of } J \text{ wrt. } \theta$

- Doesn't work with dropout.

$$J \quad \underline{\text{keep-prob} = 1.0}$$

- Run at random initialization; perhaps again after some training.

$$\underline{w, b \approx 0}$$